

Software Defined Networking

Lab Work 2 Solution



TECHNISCHE
UNIVERSITÄT
DARMSTADT

Jeremias Blendin, Leonhard Nobach, Christian Koch,
Julius Rückert, Matthias Wichtlhuber



PS - Peer-to-Peer Systems Engineering Lab
Dept. of Electrical Engineering and Information Technology
Technische Universität Darmstadt
Rundeturmstr. 12, D-64283 Darmstadt, Germany
<http://www.ps.tu-darmstadt.de/>



INTRODUCTION TO OpenFlow Controller / RYU

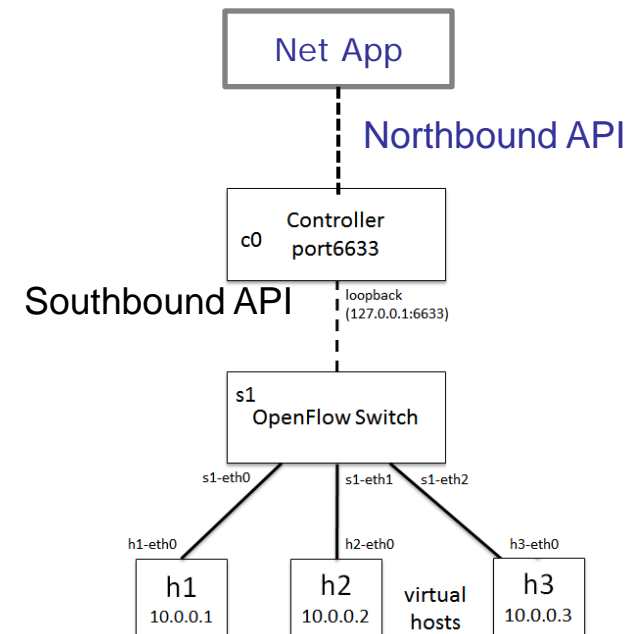
Short Recap

❖ Previously we manually added rules in the switch

➤ `$ dpctl add-flow tcp:127.0.0.1:6634\
in_port=1,idle_timeout=0,actions=output:2`

❖ This should be done automatically

- Task of a Network Application (App)
- E.g. a simple switching App



[1] <http://sdnhub.org/resources/useful-mininet-setups/>

Installing Ryu

❖ Reboot your existing Mininet VM and enter:

- `$ sudo -s`
- `$ apt-get install python-eventlet python-routes
python-webob python-paramiko python-pip python-dev
libxml2-dev libxslt-dev zlib1g-dev`
- `$ pip install ryu`
- `$ mn -c`

Run a simple_switch

❖ Enter:

- `$ mn --topo single,3 --mac --arp --switch ovsk\ --controller=remote,ip=127.0.0.1`
- `$ h1 ping h2 -> timeout`

❖ Open a second terminal and connect to the VM

➤ Execute

- `$ ryu-manager ryu.app.simple_switch --verbose`
- `h1 ping h2 ➡ success`
- Investigate the OpenFlow rules in switch s1
 - New tool: `ovs-ofctl`
 - `$ ovs-ofctl dump-flows s1`

Understand how it works

- ❖ A step-by-step explanation can be found here
 - http://osrg.github.io/ryu-book/en/html/switching_hub.html
 - Read it carefully!
- ❖ Other resources like books and tutorials available
 - E.g. <http://books.google.de/books?id=JC3rAgAAQBAJ>

Task 1: Packet Replication 1/2

- ❖ Modify the *simple_switch.py* in a way that all received ICMP request packets are sent through the two other *out_ports* of the switch. The packet should not be sent back to the port from where it originated.
 - The basis for the task is the Ryu application *simple_switch.py* and OF 1.0:
https://github.com/osrg/ryu/blob/master/ryu/app/simple_switch.py
 - A ping request from h1 to h2 should result in a ping reply to h1 from h2 and h3. As a result, h1 receives more packets than it has sent.
 - It is sufficient for the solution to work for in the example network with 3 hosts
 - Mininet provides a fixed mapping between OpenFlow port numbers, MAC, and IP addresses. This information should be used for implementation.
 - Carefully think about what actions need to be applied to the ICMP packets
 - Have a look at the respective standards documents:
 - OpenFlow Switch Specification 1.0.0 & Errata
<https://www.opennetworking.org/sdn-resources/technical-library>

Task 1: Packet Replication 2/2



❖ Debugging

- How to open a third terminal and connect it to one of the hosts?
 - Use xterm if you have a GUI installed
 - Open a second ssh session to the mininet VM
 - In Mininet run: `mininet> py h3.pid -> 3013`
 - Attach to h3 by running `$ sudo mnexec -a 3013 bash`
 - Verify by running `$ ip a ->` you should see the interfaces of h3 only
- Run tcpdump on h3
 - Open a terminal on h3
 - `$ tcpdump -eUvi h3-eth0`
 - `mininet> h1 ping h3`
 - With the packet duplication code in place
 - `mininet> h1 ping h2`
 - A packet copy of the ICMP request to h2 should be visible in the tcpdump output
 - Why does h3 not send an ICMP reply to the ICMP request?

Lab 2 Task 1 Solution (1)

❖ Some confusion/uncertainty regarding the task

- We will stick to this line in the description: "A ping from h1 to h2 should result in a ping from h1 to h2 and h3. As a result, h1 receives more packets than it has sent."

```
mininet> h1 ping h2
64 bytes from 10.0.0.2: icmp_seq=343 ttl=64 time=0.076 ms
64 bytes from 10.0.0.2: icmp_seq=343 ttl=64 time=0.079 ms (DUP!)
```

- Verify using tcpdump

```
:~$ sudo tcpdump -ei s1-eth1
52 00:00:00:00:00:01 > 00:00:00:00:00:02, ethertype IPv4 (0x0800), length 98: 10.0.0.1 > 10.0.0.2:
ICMP echo request, id 5168, seq 48, length 64
92 00:00:00:00:00:02 > 00:00:00:00:00:01, ethertype IPv4 (0x0800), length 98: 10.0.0.2 > 10.0.0.1:
ICMP echo reply, id 5168, seq 48, length 64
95 00:00:00:00:00:03 > 00:00:00:00:00:01, ethertype IPv4 (0x0800), length 98: 10.0.0.3 > 10.0.0.1:
ICMP echo reply, id 5168, seq 48, length 64
```

Lab 2 Task 1 Sample Solution (1)

- ❖ See the file „simple_switch_duplicate.py“ for a complete example solution
- ❖ The network is static and known: Add a host list

```
# Code by Patrick Welzel and Mahshid Okhovatzadeh
self.hosts_for_omniping = [
    {'port': 1,
     'ip': '10.0.0.1',
     'mac': '00:00:00:00:00:1'},
    {'port': 2,
     'ip': '10.0.0.2',
     'mac': '00:00:00:00:00:2'},
    {'port': 3,
     'ip': '10.0.0.3',
     'mac': '00:00:00:00:00:3'},
]
```

Lab 2 Task 1 Sample Solution (2)



TECHNISCHE
UNIVERSITÄT
DARMSTADT

❖ Add additional output rules

```
Adarsh Chikkaballapur Umashankar
Bhargava Narasipura
actions=[]
# install a flow to avoid packet_in next time
if out_port != ofproto.OFPP_FLOOD:
    for host in self.hosts_for_omniping:
        if msg.in_port != host['port']:
            if out_port != host['port']:
                actions.append(datapath.ofproto_parser.OFActionSetDlDst(haddr_to_bin(host['mac'])))
                actions.append(datapath.ofproto_parser.OFActionSetNwDst(ofctl.ipv4_to_int(host['ip'])))
                actions.append(datapath.ofproto_parser.OFActionOutput(host['port']))
            self.logger.info("actions is %s ", actions)
            self.add_flow(datapath, msg.in_port, dst, actions)
```

❖ OpenFlow rules

```
mininet@mininet-vm:~$ sudo ovs-ofctl dump-flows s1
NXST_FLOW reply (xid=0x4):
  cookie=0x0, duration=41.308s, table=0, n_packets=32, n_bytes=3136, idle_timeout=20, idle_age=9,
  icmp,in_port=2 actions=output:1,mod_dl_dst:00:00:00:00:00:03,mod_nw_dst:10.0.0.3,output:3
  cookie=0x0, duration=41.311s, table=0, n_packets=32, n_bytes=3136, idle_timeout=20, idle_age=9,
  icmp,in_port=1 actions=output:2,mod_dl_dst:00:00:00:00:00:03,mod_nw_dst:10.0.0.3,output:3
  cookie=0x0, duration=41.308s, table=0, n_packets=32, n_bytes=3136, idle_timeout=20, idle_age=9,
  icmp,in_port=3 actions=output:1,mod_dl_dst:00:00:00:00:00:02,mod_nw_dst:10.0.0.2,output:2
```



Solution: Terminal Output

Terminal1

```
mininet@mininet-vm:~$ sudo mn --topo single,3 --mac --arp --switch ovsk --  
controller=remote,ip=127.0.0.1
```

Terminal2

```
mininet@mininet-vm:~$ sudo ryu-manager ./simple_switch_duplicate.py  
loading app ./simple_switch_duplicate.py  
loading app ryu.controller.ofp_handler  
instantiating app ./simple_switch_duplicate.py of SimpleSwitch app  
ryu.controller.ofp_handler of OFPHandler
```

Terminal1

```
mininet> h1 ping h2  
PING 10.0.0.2 (10.0.0.2) 56(84) bytes of data.  
64 bytes from 10.0.0.2: icmp_seq=1 ttl=64 time=2.99 ms  
64 bytes from 10.0.0.2: icmp_seq=1 ttl=64 time=3.14 ms (DUP!)
```